

BUILDING A MINI HPC CLUSTER

FROM SCRATCH ON ROCKY LINUX

A Practical Guide to Cluster Setup, Workload Management,
Parallel Programming and HPC Infrastructure



ROCKY LINUX

Stable. Secure.
Enterprise Ready.



SLURM

Workload
Manager



OPENMPI

Parallel
Runtime



LDAP

Centralized
Authentication



NFS

Shared
Storage



SPACK

Software
Management

LEARN. BUILD. UNDERSTAND. OPTIMIZE.



Understand the
HPC Stack



Build Your
Environment



Run Parallel.
Scale Efficiently.



Think Scalable.
Optimize Better.

ClusterSetup

Cisco Ramon

[2026-03-14 Sat 00:35]

Contents

1	Abstract	9
2	Preface	9
3	About the Author	11
4	Introduction	12
5	Cluster Architecture	13
6	Node Details	14
6.1	Node Role	14
6.2	Networking Setup	15
6.3	Static IP Addressing	15
7	Resource Allocation	16
7.1	Master Node	16
7.2	Login Node	16
7.3	Compute Nodes	16
8	Master Node Setup	16
8.1	Install VirtualBox on Manjaro	17
8.2	Download Rocky Linux	17
8.3	Create Master VM	18
8.4	Network Configuration	18
8.5	Rocky Linux Installation	19
8.6	Initial System Verification	20
8.7	Configure Static Internal Network	20

8.8	Configure Internal Interface	20
8.9	Restart Networking	21
8.10	Verify Networking	21
8.11	SSH Access from Host	21
8.12	System Update and Packages	21
8.13	Timezone Configuration	22
8.14	Configure /etc/hosts	22
8.15	Disable Firewall (Temporary)	22
8.16	SELinux Configuration	22
8.17	Verify Hostname	23
9	Cloning Nodes	23
9.1	Create Snapshot	23
9.2	Clone Nodes	24
9.3	Clone Names	24
9.4	Configure Port Forwarding	24
9.5	Machine ID Regeneration	24
9.6	Remove Existing Machine ID	25
9.7	Generate New Machine ID	25
9.8	Verify Machine ID	25
9.9	Configure Hostname	25
9.10	Configure Static Internal IPs	26
9.11	Restart Networking	26
9.12	Verify Networking	27
9.13	Verify Routing	27
9.14	Verify Connectivity	27
9.15	SSH Verification	27
10	Shared Infrastructure Services	28
10.1	Cluster Role Architecture	28
10.2	Verify User Consistency	29
10.3	Passwordless SSH Configuration	30
10.4	Generate SSH Key on Login Node	30
10.5	Copy SSH Keys to Cluster Nodes	30
10.6	Verify Passwordless SSH	31
10.7	HPC Concept — Why Passwordless SSH Matters	32
10.8	NFS Shared Storage Setup	32
10.9	Install NFS Packages	32
10.10	Configure NFS Server on Master	33
10.11	Enable NFS Services	33

10.12	Configure NFS Exports	33
10.13	Export Option Explanation	33
10.14	Apply Exports	33
10.15	Verify Exports	33
10.16	Configure NFS Clients	34
10.17	Backup Existing Home Directory	34
10.18	Create New Mount Point	34
10.19	Mount Shared NFS Home	34
10.20	Verify Mount	34
10.21	Verify Shared Filesystem	35
10.22	Configure Persistent NFS Mounts	35
10.23	Configure /etc/fstab	35
10.24	Verify fstab Before Reboot	35
10.25	Reboot Validation Procedure	36
10.26	Reboot Order	36
10.27	Verify Persistent Mounts After Reboot	36
10.28	Verify Active NFS Mount	37
10.29	Verify Shared Files After Reboot	37
11	Munge Authentication + Slurm Scheduler	37
11.1	HPC Scheduler Architecture	38
11.2	IMPORTANT Rocky Linux 10 Issue	39
11.3	Install EPEL Repository	39
11.4	Enable CRB Repository	39
11.5	Install Munge + Slurm Packages	40
11.6	Verify Installed Packages	40
11.7	Verify System Users	40
11.8	IMPORTANT Issue Encountered — Missing slurm User .	41
11.9	Create Slurm Group	41
11.10	Create Slurm User	41
11.11	Verify Slurm User	41
11.12	IMPORTANT HPC Security Concept	41
11.13	Configure Munge Authentication	42
11.14	Generate Munge Key on Master	42
11.15	Verify Munge Key	42
11.16	IMPORTANT Mistake Encountered — Munge Key Gen- erated on All Nodes	42
11.17	Correct Munge Key Distribution	43
11.18	IMPORTANT SSH Root Login Issue	43
11.19	Copy Munge Key to Login Node	43

11.20	Copy Munge Key to Compute01	43
11.21	Copy Munge Key to Compute02	43
11.22	Copy Munge Key to Compute03	44
11.23	Remove Temporary Key Copy	44
11.24	Install Munge Key on Client Nodes	44
11.25	Stop Munge Service	44
11.26	Install Key	44
11.27	Fix Ownership	44
11.28	Fix Permissions	44
11.29	Start Munge	44
11.30	IMPORTANT Issue Encountered — Munge Not Running on Master	45
11.31	Verify Munge Service	45
11.32	Verify Munge Authentication	45
11.33	Test Compute02	45
11.34	Test Compute03	45
11.35	Test Login Node	46
11.36	HPC Concepts Learned	46
11.37	Configure Slurm Scheduler	46
11.38	IMPORTANT Issue Encountered — Default localhost Con- figuration	46
11.39	Create Slurm Configuration	47
11.40	Configuration Explanation	48
11.41	Create Slurm Runtime Directories	48
11.42	Set Ownership	48
11.43	Create slurmd Directories on Compute Nodes	48
11.44	IMPORTANT sudo Over SSH Issue	49
11.45	Fix Ownership on Compute Nodes	49
11.46	Verify Runtime Directories	49
11.47	Distribute slurm.conf	50
11.48	Install Configuration on Nodes	50
11.49	Start Slurm Controller	50
11.50	Verify Controller	50
11.51	Start slurmd on Compute Nodes	51
11.52	Verify slurmd	51
11.53	Reconfigure Slurm	51
11.54	IMPORTANT Issue Encountered — Nodes in DRAIN State	51
11.55	Stop slurmd	52
11.56	Remove Old State	52
11.57	Restart slurmd	52

11.58	Resume Nodes	52
11.59	Verify Cluster Status	52
11.60	Verify Node Details	52
11.61	Test Slurm Execution	53
11.62	Multi-node Execution	53
11.63	Multi-task Multi-node Execution	53
11.64	Multi-core Single-node Execution	53
11.65	IMPORTANT Slurm Scheduling Concept	54
12	Shared Software Filesystem (/apps)	54
12.1	Why Separate /apps Filesystem?	55
12.2	IMPORTANT Storage Architecture Concept	55
12.3	Add New Virtual Disk to Master	56
12.4	Verify New Disk Detection	57
12.5	Create GPT Partition Table	57
12.6	Verify Partition	58
12.7	Create XFS Filesystem	58
12.8	Create Mount Point	58
12.9	Get Filesystem UUID	58
12.10	Configure Persistent Mount	58
12.11	Mount Filesystem	59
12.12	Verify Mount	59
12.13	Configure Ownership	59
12.14	Create Shared HPC Software Hierarchy	59
12.15	HPC Software Hierarchy Concept	60
12.16	Export /apps via NFS	60
12.17	Apply NFS Exports	60
12.18	Verify Exports	60
12.19	Create Mount Points on Client Nodes	60
12.20	Configure Persistent NFS Mounts	61
12.21	Mount Shared /apps Filesystem	61
12.22	Verify Shared Mount	61
12.23	Verify Shared Filesystem Functionality	61
12.24	HPC Architecture After /apps Setup	62
13	Spack Software Stack + MPI Runtime Validation	62
13.1	IMPORTANT HPC Software Stack Concept	63
13.2	IMPORTANT Shared Filesystem Concept	64
13.3	Install Development Tools	64
13.4	Move into Shared Application Filesystem	65

13.5	Clone Spack into Shared Filesystem	65
13.6	Enter Spack Directory	66
13.7	Initialize Spack Environment	66
13.8	IMPORTANT Shell Environment Concept	67
13.9	Make Spack Permanent	67
13.10	Clone Additional Spack Repository	68
13.11	Configure Spack Repository	68
13.12	Verify Spack Functionality	68
13.13	IMPORTANT HPC Software Stack Concept	69
13.14	IMPORTANT Spack Build Concept	69
13.15	Detect Available Compilers	70
13.16	Install GCC Using Spack	70
13.17	IMPORTANT Multiple Compiler Version Concept	70
13.18	Verify Installed Software	71
13.19	Load Installed GCC Compiler	71
13.20	IMPORTANT Runtime Environment Concept	72
13.21	Register Newly Installed Compiler	72
13.22	Install OpenMPI Using Spack	72
13.23	IMPORTANT MPI Compiler Dependency Concept	73
13.24	Verify Installed Software Tree	73
13.25	IMPORTANT Spack Spec Concept	74
13.26	IMPORTANT Hash-Based Software Loading	74
13.27	Load OpenMPI Environment	75
13.28	IMPORTANT MPI Wrapper Concept	75
13.29	Verify Loaded Software	75
13.30	Create MPI Performance Validation Program	76
13.31	IMPORTANT MPI Reduction Concept	79
13.32	IMPORTANT Parallel Execution Concept	80
13.33	Load Software Before Compilation	80
13.34	Compile MPI Program	81
13.35	IMPORTANT MPI Compilation Concept	81
13.36	Interactive MPI Runtime Validation	82
13.37	IMPORTANT Shared Filesystem Advantage	82
13.38	Create Slurm Batch Script	83
13.39	IMPORTANT Slurm Script Explanation	83
13.40	IMPORTANT Runtime Environment Section	84
13.41	IMPORTANT Slurm Runtime Variable	84
13.42	Submit MPI Job	85
13.43	Monitor Jobs	85
13.44	IMPORTANT Job States	85

13.45	View Detailed Job Information	86
13.46	Check Cluster Status	86
13.47	IMPORTANT Node States	86
13.48	Check Output Files	87
13.49	IMPORTANT HPC Performance Metrics	87
13.50	Cancel Jobs	88
13.51	Interactive Slurm Execution	88
13.52	IMPORTANT HPC Workflow	89
13.53	IMPORTANT Production HPC Concept	89
14	Centralized Identity Management using OpenLDAP + nslcd	90
14.1	Important Architecture	91
14.2	Important Linux Authentication Concepts	91
14.3	Install OpenLDAP Server on Master Node	91
14.4	Enable LDAP Server	92
14.5	Generate LDAP Admin Password Hash	92
14.6	Configure LDAP Database	93
14.7	Important LDAP Concepts	94
14.8	Apply LDAP Database Configuration	94
14.9	Import LDAP Schemas	95
14.10	Create LDAP Base Tree	95
14.11	Verify LDAP Tree	96
14.12	Create LDAP Group	97
14.13	Create LDAP User Password Hash	97
14.14	Create LDAP User	98
14.15	Important LDAP User Attributes	99
14.16	Verify LDAP User	99
14.17	Configure LDAP Client on Login Node	99
14.18	Configure nslcd	100
14.19	Enable nslcd	100
14.20	Verify LDAP Connectivity	101
14.21	Configure NSS	101
14.22	Restart nslcd	102
14.23	Verify LDAP User Resolution	102
14.24	NOTE (Do this for master too)	102
14.25	Install Automatic Home Directory Support	102
14.26	Configure PAM Authentication	103
14.27	Configure Second PAM File	104
14.28	Restart Services	104

14.29	Verify LDAP Authentication	104
14.30	Verify User Home	105
14.31	Troubleshooting	105
14.32	Current Cluster State After LDAP Setup	110
14.33	Important Debugging Commands	111
14.34	Important Architecture Understanding	111
15	Automated LDAP User Creation Script	112
15.1	Script Location	112
15.2	User Creation Script	113
15.3	Make Script Executable	115
15.4	Understanding Script Arguments	116
15.5	Default Password Format	117
15.6	What the Script Automatically Does	117
15.7	Example User Creation	117
15.8	Verify User Creation	118
15.9	Verify Login	118
15.10	Important UID/GID Notes	119
15.11	Important HPC Identity Concept	119
16	Automated LDAP User Deletion Script	120
16.1	Script Location	120
16.2	User Deletion Script	121
16.3	Make Script Executable	123
16.4	Understanding Script Arguments	123
16.5	What the Script Automatically Does	124
16.6	Example User Deletion	124
16.7	Important Group Deletion Notes	125
16.8	Important Home Directory Notes	125
16.9	Verify User Removal	126
16.10	Important HPC User Management Concept	126
17	Automated LDAP Password Change Script	127
17.1	Script Location	127
17.2	Password Change Script	128
17.3	Make Script Executable	130
17.4	Understanding Script Arguments	130
17.5	What the Script Automatically Does	130
17.6	Example Password Change	130
17.7	Verify Password Change	131

17.8	Important Password Management Notes	132
------	---	-----

1 Abstract

This documentation presents the complete setup and configuration of a lightweight High Performance Computing (HPC) cluster environment using Rocky Linux and VirtualBox. The guide covers practical implementation of core HPC infrastructure components including shared storage, Slurm workload management, LDAP-based centralized authentication, MPI runtime environment, and centralized software stack management using Spack.

The cluster environment supports multinode job execution, centralized software distribution, distributed MPI applications, and scheduler-aware runtime execution. The setup follows real-world HPC concepts while remaining suitable for learning, experimentation, and practical hands-on training.

2 Preface

This documentation was prepared with the assistance of AI tools for:

- structuring
- formatting
- workflow organization
- command documentation
- concept explanations

However, the complete HPC cluster setup, configuration, testing, and validation were performed manually by me in a fully working VirtualBox based cluster environment.

All commands, scripts, workflows, and configurations documented in this guide were tested successfully during the setup process, including:

- cluster networking
- shared storage
- NFS configuration
- Slurm workload manager

- LDAP authentication
- centralized user management
- Spack software stack
- MPI runtime environment
- multinode execution

This guide is intended for:

- students
- beginners
- Linux learners
- HPC enthusiasts
- system administrators
- researchers

who want practical hands-on experience in:

- HPC infrastructure setup
- cluster administration
- workload scheduling
- centralized authentication
- parallel programming
- MPI execution
- software stack management

The goal of this documentation is not only to provide commands, but also to explain the practical concepts behind real HPC environments in a simple and beginner-friendly manner.

3 About the Author

Hi, I'm Abhishek Raj

I work in the field of High Performance Computing (HPC), Linux systems, parallel programming, and scientific computing with experience in HPC application optimization and distributed computing environments, including the PARAM series of supercomputers.

My primary areas of interest include:

- HPC clusters and infrastructure
- Linux systems and administration
- MPI, OpenMP, CUDA, and OpenACC
- scientific and research workloads
- distributed systems
- software optimization and performance tuning
- automation and reproducible environments

This project was created to gain a deeper understanding of how HPC clusters are designed and managed internally. The objective was not only to build a working cluster environment from scratch, but also to better understand the complete software and infrastructure stack behind parallel computing systems, which ultimately helps in writing and optimizing efficient HPC applications.

GitHub Profile:

<https://github.com/CISSSCO/>

LinkedIn :

<https://www.linkedin.com/in/abhi581b>

Project Repository:

<https://github.com/CISSSCO/hpccfs>

Hosted Documentation:

<https://cisssco.github.io/hpccfs>

Source codes, scripts, configuration files, and examples used in this guide are available in the project repository.

For queries, suggestions, or discussions:

<https://ciscoramon.netlify.app/>

If you encounter any issue in the documentation, setup process, or configuration steps, feel free to create an issue in the GitHub repository.

4 Introduction

High Performance Computing (HPC) clusters are distributed computing systems designed to execute computational workloads across multiple machines simultaneously. Instead of relying on a single system, HPC clusters combine multiple compute nodes to provide higher computational capacity, scalability, and parallel execution capability.

Modern HPC environments are widely used in:

- scientific simulations
- engineering workloads
- computational chemistry
- artificial intelligence
- machine learning
- weather modeling
- large-scale data analysis

A typical HPC cluster contains several important components working together:

Component	Purpose
Login Node	User access and job submission
Compute Nodes	Execute computational workloads
Shared Storage	Shared user data and software
Scheduler	Resource allocation and job management
MPI Runtime	Distributed parallel execution
Authentication System	Centralized user management

This guide focuses on building these components step-by-step in a small virtualized environment using:

- Rocky Linux
- VirtualBox
- Slurm
- OpenMPI

- LDAP
- NFS
- Spack

Although the setup is lightweight and designed for learning purposes, the overall architecture closely follows concepts used in production HPC centers and supercomputers.


Throughout this documentation, readers will gradually configure:

- cluster networking
- shared storage
- centralized authentication
- workload scheduling
- distributed software stacks
- MPI-based parallel execution
- scheduler-aware job execution

The final environment provides a fully working multinode HPC cluster capable of running distributed parallel applications across multiple compute nodes.

5 Cluster Architecture

The cluster consists of five virtual machines connected using an internal network.



master	10.10.10.1
login	10.10.10.2
compute01	10.10.10.11
compute02	10.10.10.12
compute03	10.10.10.13

6 Node Details

6.1 Node Role

1. Master Node

The master node manages the overall cluster.

Main responsibilities:

- Slurm controller
- NFS server
- LDAP server
- Monitoring tools
- Shared software stack
- Cluster configuration

2. Login Node

The login node is the access point for users.

Typical tasks performed here:

- SSH login
- Code compilation
- Job submission
- Accessing shared storage

3. Compute Nodes

Compute nodes run the actual workloads.

These nodes are mainly used for:

- MPI jobs

- OpenMP jobs
- Batch workloads
- Parallel execution

6.2 Networking Setup

Two network adapters are used for each VM.

1. NAT Adapter

Used for:

- Internet access
- Package installation
- SSH access from host system

2. Internal Network

Used for communication between cluster nodes.

Network name:

`hpcnet`

This network is isolated and only used inside the cluster.

6.3 Static IP Addressing

Each node uses a fixed IP address.

VM Name	Hostname	Internal IP	SSH Port
master	master.local	10.10.10.1	2221
login	login.local	10.10.10.2	2222
compute01	compute01.local	10.10.10.11	2223
compute02	compute02.local	10.10.10.12	2224
compute03	compute03.local	10.10.10.13	2225

Example SSH access:

```
ssh -p 2221 hpcuser@localhost
```


7 Resource Allocation

The following configuration is enough for a lightweight HPC lab. You can increase resources later depending on your host system.

7.1 Master Node

Resource	Value
RAM	2 GB
CPU	2
Disk	30 GB

7.2 Login Node

Resource	Value
RAM	2 GB
CPU	2
Disk	25 GB

7.3 Compute Nodes

Resource	Value
RAM	2 GB
CPU	2
Disk	20 GB

8 Master Node Setup

This section covers the complete setup of the master node. The master node acts as the central management node of the cluster.

Main responsibilities:

- Cluster management
- Shared storage
- Slurm controller
- Monitoring services
- User authentication
- Software management

8.1 Install VirtualBox on Manjaro

I'm using Manjaro (Arch) as my base system.

1. Update packages

```
sudo pacman -Syu
```

2. Install VirtualBox and host modules

```
sudo pacman -S virtualbox virtualbox-host-modules-arch
```

3. Install VirtualBox extension pack

```
sudo pacman -S virtualbox-ext-oracle
```

4. Load VirtualBox kernel module

```
sudo modprobe vboxdrv
```

5. Add current user to VirtualBox group

```
sudo usermod -aG vboxusers $USER
```

6. Reboot the system

```
sudo reboot
```

8.2 Download Rocky Linux

Download Rocky Linux Minimal ISO:

<https://rockylinux.org/download>

Choose:

- x86_64
- Minimal ISO

The minimal installation keeps the environment lightweight and avoids unnecessary packages.

8.3 Create Master VM

Open VirtualBox and create a new virtual machine.

1. VM Name

`master`

2. Type

`Linux`
`Red Hat (64-bit)`

3. Hardware

Setting	Value
RAM	2046 MB
CPU	2
Disk	30 GB

4. Disk Type

`VDI`
`Dynamically Allocated`

This configuration is enough for a lightweight HPC lab environment.

8.4 Network Configuration

1. Adapter 1

`Attached to: NAT`

`Used for:`

- Internet access
- Package installation
- SSH access from host machine

2. NAT Port Forwarding

Open:

Settings -> Network -> Adapter 1 -> Advanced -> Port Forwarding

Name	Protocol	Host Port	Guest Port
ssh	TCP	2221	22

3. Adapter 2

Attached to: Internal Network

Internal Network Name:

hpcnet

IMPORTANT: All nodes MUST use the same internal network name.
This network is used for communication between cluster nodes.

8.5 Rocky Linux Installation

Choose:

Minimal Install

1. Network and Hostname

Enable BOTH network adapters.

Set hostname:

master.local

2. User Creation

Create an administrative user during installation.

Field	Value
Username	hpcadmin
Administrator	YES

After installation, log into the VM.

8.6 Initial System Verification

1. Check Interfaces

```
ip addr
```

Expected:

- NAT interface should receive:
 - 10.0.2.15
- Internal interface should exist but may not yet have static IP

8.7 Configure Static Internal Network

1. Check Device Names

```
nmcli device status
```

Expected:

Interface	Purpose
enp0s3	NAT
enp0s8	Internal Network

8.8 Configure Internal Interface

Launch NetworkManager TUI:

```
sudo nmtui
```

Navigate:

Edit a connection

Select:

enp0s8

Set:

Field	Value
IPv4 Method	Manual
Address	10.10.10.1/24
Gateway	blank
DNS	blank

DO NOT modify NAT interface.

8.9 Restart Networking

```
sudo systemctl restart NetworkManager
```

8.10 Verify Networking

```
ip addr
```

Expected:

Interface	IP
NAT	10.0.2.x
Internal	10.10.10.1

8.11 SSH Access from Host

From host machine:

```
ssh -p 2221 hpcadmin@localhost
```

If the connection works, port forwarding is configured correctly.

8.12 System Update and Packages

1. Update System

```
sudo dnf update -y
```

2. Install Base Packages

```
sudo dnf install -y \  
vim wget curl git \  
net-tools bind-utils \  
htop tmux tree \  
gcc gcc-c++ gcc-gfortran \  
make cmake \  
python3 \  
openssh-server
```

These packages provide:

- Development tools
- Network utilities
- Monitoring utilities
- Basic administration tools

8.13 Timezone Configuration

Check timezone:

```
timedatectl
```

Set timezone if needed:

```
sudo timedatectl set-timezone Asia/Kolkata
```

8.14 Configure /etc/hosts

Edit:

```
sudo vim /etc/hosts
```

Add:

```
10.10.10.1    master master.local
10.10.10.2    login login.local
10.10.10.11   compute01 compute01.local
10.10.10.12   compute02 compute02.local
10.10.10.13   compute03 compute03.local
```

This allows hostname-based communication between nodes.

8.15 Disable Firewall (Temporary)

For initial setup and testing:

```
sudo systemctl disable --now firewalld
```

NOTE:

Firewall rules can be configured properly later after cluster setup is complete.

8.16 SELinux Configuration

Check status:

```
getenforce
```

Temporarily disable enforcing:

```
sudo setenforce 0
```

Persist configuration:

```
sudo vim /etc/selinux/config
```

Modify:

```
SELINUX=permissive
```

This helps avoid permission-related issues during initial cluster setup.

8.17 Verify Hostname

```
hostnamectl
```

Expected:

```
Static hostname: master.local
```

If hostname is not updated(showing previous), reboot the machine.

At this point, the master node setup is complete and ready for cloning or further cluster configuration.

9 Cloning Nodes

Purpose:

- avoid reinstalling all nodes manually
- simulate real HPC image provisioning workflow

9.1 Create Snapshot

Power off VM:

```
sudo poweroff
```

Create snapshot in VirtualBox:

```
clean-base-master
```


9.2 Clone Nodes

Create FULL CLONES.

IMPORTANT: Enable:

Reinitialize the MAC address of all network cards

9.3 Clone Names

VM Name
login
compute01
compute02
compute03

9.4 Configure Port Forwarding

1. Login Node

Host Port	Guest Port
2222	22

2. Compute01

Host Port	Guest Port
2223	22

3. Compute02

Host Port	Guest Port
2224	22

4. Compute03

Host Port	Guest Port
2225	22

9.5 Machine ID Regeneration

IMPORTANT: Each cloned Linux machine must have unique machine-id.

9.6 Remove Existing Machine ID

```
sudo rm -f /etc/machine-id
```

9.7 Generate New Machine ID

```
sudo systemd-machine-id-setup
```

9.8 Verify Machine ID

```
cat /etc/machine-id
```

Ensure every node has unique machine-id.

9.9 Configure Hostname

1. Login Node

Edit hostname:

```
sudo vim /etc/hostname
```

Set:

```
login.local
```

Apply:

```
sudo hostnamectl set-hostname login.local
```

2. Compute01

Set hostname:

```
compute01.local
```

Apply:

```
sudo hostnamectl set-hostname compute01.local
```

3. Compute02

```
compute02.local
```

```
sudo hostnamectl set-hostname compute02.local
```

4. Compute03

```
compute03.local
```

```
sudo hostnamectl set-hostname compute03.local
```

9.10 Configure Static Internal IPs

Launch:

```
sudo nmtui
```

Edit internal adapter:

```
enp0s8
```

1. Login Node

```
10.10.10.2/24
```

2. Compute01

```
10.10.10.11/24
```

3. Compute02

```
10.10.10.12/24
```

4. Compute03

```
10.10.10.13/24
```

9.11 Restart Networking

```
sudo systemctl restart NetworkManager
```

9.12 Verify Networking

```
ip addr
```

9.13 Verify Routing

```
ip route
```

Expected:

```
default via 10.0.2.2 dev enp0s3  
10.10.10.0/24 dev enp0s8
```

9.14 Verify Connectivity

1. Ping Between Nodes

```
ping master
```

```
ping login
```

```
ping compute01
```

9.15 SSH Verification

From host:

1. Master

```
ssh -p 2221 hpcadmin@localhost
```

2. Login

```
ssh -p 2222 hpcadmin@localhost
```

3. Compute01

```
ssh -p 2223 hpcadmin@localhost
```

4. Compute02

```
ssh -p 2224 hpcadmin@localhost
```

5. Compute03

```
ssh -p 2225 hpcadmin@localhost
```

10 Shared Infrastructure Services

Goal:

Transform the cluster into a realistic HPC-style shared environment.

This phase includes:

- passwordless SSH
- shared home directories
- NFS server/client configuration
- persistent mounts
- multi-node shared filesystem verification

By the end of this phase:

- all nodes share same /home
- users can access same files from all nodes
- cluster behaves like a real HPC environment

10.1 Cluster Role Architecture

1. Master Node

Purpose:

- cluster management
- NFS server
- future Slurm controller
- future Munge authority
- future LDAP server

Services:

Service	Purpose
nfs-server	shared storage
future slurmd	scheduler controller
future munge	authentication

2. Login Node

Purpose:

- user login
- compilation
- job submission

Users should:

- SSH into login node
- compile programs
- submit jobs using Slurm

3. Compute Nodes

Purpose:

- execute jobs only

Future services:

Service	Purpose
slurmd	compute daemon
munge	authentication

10.2 Verify User Consistency

IMPORTANT:

NFS shared storage requires same UID/GID across all nodes.

Verify user identity on all nodes:

```
id hpcadmin
```

Expected:

```
uid=1000(hpcadmin) gid=1000(hpcadmin)
```

All nodes must match.

10.3 Passwordless SSH Configuration

Purpose:

- MPI communication
- Slurm execution
- cluster automation
- remote management

SSH trust is generated ONLY from login node.
This mimics real HPC architecture:

User -> Login Node -> Compute Nodes

10.4 Generate SSH Key on Login Node

SSH into login node from host:

```
ssh -p 2222 hpcadmin@localhost
```

Generate SSH key:

```
ssh-keygen
```

Press Enter for all prompts.
Generated files:

```
~/.ssh/id_rsa  
~/.ssh/id_rsa.pub
```

10.5 Copy SSH Keys to Cluster Nodes

1. Master

```
ssh-copy-id hpcadmin@master
```

2. Compute01

```
ssh-copy-id hpcadmin@compute01
```

3. Compute02

```
ssh-copy-id hpcadmin@compute02
```

4. Compute03

```
ssh-copy-id hpcadmin@compute03
```

10.6 Verify Passwordless SSH

1. Test Master

```
ssh master hostname
```

Expected:

```
master.local
```

2. Test Compute01

```
ssh compute01 hostname
```

Expected:

```
compute01.local
```

3. Test Compute02

```
ssh compute02 hostname
```

4. Test Compute03

```
ssh compute03 hostname
```

No password prompts should appear.

10.7 HPC Concept — Why Passwordless SSH Matters

Passwordless SSH becomes essential for:

Feature	Purpose
MPI	remote process launch
Slurm	distributed execution
automation	orchestration
monitoring	remote administration

10.8 NFS Shared Storage Setup

Purpose:

Provide shared home directories across all cluster nodes.

This is a core HPC architecture requirement.

Without shared storage:

- source code differs per node
- binaries differ
- scripts unavailable across nodes
- MPI execution becomes difficult

Real HPC systems commonly use:

- NFS
- Lustre
- BeeGFS
- GPFS

This lab begins with NFS.

10.9 Install NFS Packages

Install on ALL nodes:

```
sudo dnf install -y nfs-utils
```

10.10 Configure NFS Server on Master

SSH into master node.

10.11 Enable NFS Services

```
sudo systemctl enable --now nfs-server
```

Verify:

```
systemctl status nfs-server
```

Expected:

```
active (running)
```

10.12 Configure NFS Exports

Edit exports file:

```
sudo vim /etc/exports
```

Add:

```
/home 10.10.10.0/24(rw,sync,no_root_squash)
```

10.13 Export Option Explanation

Option	Meaning
rw	read/write access
sync	synchronous writes
no_root_squash	preserve root privileges

10.14 Apply Exports

```
sudo exportfs -rav
```

10.15 Verify Exports

```
sudo exportfs -v
```

Expected export:

```
/home 10.10.10.0/24(...)
```

10.16 Configure NFS Clients

Perform on:

- login
- compute01
- compute02
- compute03

10.17 Backup Existing Home Directory

```
sudo mv /home /home.local
```

Purpose:

Preserve original local home directory.

10.18 Create New Mount Point

```
sudo mkdir /home
```

10.19 Mount Shared NFS Home

```
sudo mount master:/home /home
```

10.20 Verify Mount

```
df -h
```

Expected:

```
master:/home
```

mounted on:

```
/home
```

10.21 Verify Shared Filesystem

From login node:

```
touch ~/shared_test
```

From compute01:

```
ls ~
```

Expected:

```
shared_test
```

This confirms successful shared filesystem operation.

10.22 Configure Persistent NFS Mounts

Without this configuration:

- NFS mount disappears after reboot

10.23 Configure /etc/fstab

On ALL client nodes:

```
sudo vim /etc/fstab
```

Add:

```
master:/home /home nfs defaults 0 0
```

10.24 Verify fstab Before Reboot

IMPORTANT:

Always validate fstab before rebooting.

Run:

```
sudo mount -a
```

If no errors appear:

- fstab configuration is correct

10.25 Reboot Validation Procedure

Purpose:

Validate persistent cluster infrastructure after reboot.

This is important real-world HPC administration practice.

10.26 Reboot Order

1. Reboot Master First

```
sudo reboot
```

Wait until:

- SSH becomes available
- NFS server active

Verify:

```
systemctl status nfs-server
```

2. Reboot Login Node

```
sudo reboot
```

3. Reboot Compute Nodes

```
sudo reboot
```

Perform for:

- compute01
- compute02
- compute03

10.27 Verify Persistent Mounts After Reboot

On login node:

```
df -h
```

Expected:

```
master:/home
```

mounted automatically.

10.28 Verify Active NFS Mount

```
mount | grep home
```

Expected:

```
master:/home on /home type nfs
```

10.29 Verify Shared Files After Reboot

```
ls ~
```

Expected:

- previously created shared files still visible

11 Munge Authentication + Slurm Scheduler

Goal:

Transform the cluster into a scheduler-managed HPC environment.

This phase includes:

- Munge authentication
- Slurm controller setup
- Compute node daemon setup
- Scheduler configuration
- Partition creation
- Multi-node execution
- Resource allocation testing

By the end of this phase:

- cluster nodes authenticate securely
- Slurm manages compute resources
- jobs execute across nodes
- scheduler topology becomes operational

11.1 HPC Scheduler Architecture

1. Master Node

Purpose:

- Slurm controller
- cluster scheduler
- resource manager
- Munge authority

Services:

Service	Purpose
munged	cluster authentication
slurmctld	scheduler controller

2. Login Node

Purpose:

- user login
- compilation
- job submission

Users should:

- login here
- compile here
- submit jobs here

3. Compute Nodes

Purpose:

- execute jobs

Services:

Service	Purpose
munged	authentication
slurmd	compute execution daemon

11.2 IMPORTANT Rocky Linux 10 Issue

Originally the cluster was created using:

Rocky Linux 10

Problem encountered:

No match for argument: slurm

Reason:

- Rocky Linux 10 HPC ecosystem still immature
- Slurm packages unavailable in standard repos
- OpenHPC EL10 support incomplete

Resolution:

- cluster rebuilt using Rocky Linux 9 Minimal

Important HPC lesson:

Production HPC clusters usually avoid bleeding-edge enterprise releases because:

- scheduler ecosystems lag behind
- HPC repos may not exist
- compatibility becomes unstable

11.3 Install EPEL Repository

Install on ALL nodes:

```
sudo dnf install -y epel-release
```

11.4 Enable CRB Repository

Required on Rocky Linux 9 and 10.

Install on ALL nodes:

```
sudo crb enable
```


11.5 Install Munge + Slurm Packages

Install on ALL nodes:

```
sudo dnf install -y \  
munge munge-libs \  
slurm slurm-slurmd slurm-slurmctld
```

11.6 Verify Installed Packages

1. Verify Munge

```
rpm -qa | grep munge
```

Expected:

```
munge  
munge-libs
```

2. Verify Slurm

```
rpm -qa | grep slurm
```

Expected:

```
slurm  
slurm-slurmd  
slurm-slurmctld
```

11.7 Verify System Users

1. Verify Munge User

```
id munge
```

2. Verify Slurm User

```
id slurm
```

11.8 IMPORTANT Issue Encountered — Missing slurm User

Problem:

```
id: 'slurm': no such user
```

Reason:

Some Slurm packages on Rocky Linux do not automatically create the slurm user.

Resolution:

Create manually on ALL nodes.

11.9 Create Slurm Group

```
sudo groupadd slurm
```

11.10 Create Slurm User

```
sudo useradd -r -g slurm -s /sbin/nologin slurm
```

Explanation:

Option	Meaning
-r	system account
-g slurm	primary group
-s /sbin/nologin	disable shell login

11.11 Verify Slurm User

```
id slurm
```

Expected:

```
uid=...
```

```
gid=...
```

11.12 IMPORTANT HPC Security Concept

Cluster daemons should not run permanently as root.

Therefore:

- Munge runs as munge user
- Slurm runs as slurm user

This is standard enterprise daemon security practice.

11.13 Configure Munge Authentication

Purpose:

Create secure trust relationship between all cluster nodes.

IMPORTANT:

All nodes MUST use the same:

```
/etc/munge/munge.key
```

This key acts as the cluster-wide shared secret.

11.14 Generate Munge Key on Master

ONLY on master node:

```
sudo /usr/sbin/create-munge-key
```

11.15 Verify Munge Key

```
ls -l /etc/munge/munge.key
```

Expected:

```
-rw-----
```

owned by:

```
munge munge
```

11.16 IMPORTANT Mistake Encountered — Munge Key Generated on All Nodes

Problem:

Munge key accidentally generated independently on all nodes.

Result:

Nodes could not authenticate each other.

Important HPC lesson:

Munge works using a:

```
shared trust model
```

All nodes MUST share identical munge.key.

11.17 Correct Munge Key Distribution

1. Copy Key Temporarily

On master:

```
sudo cp /etc/munge/munge.key /tmp/
```

2. Temporary Permission Adjustment

```
sudo chmod 644 /tmp/munge.key
```

Purpose:

Allow normal user SCP transfer.

11.18 IMPORTANT SSH Root Login Issue

Problem encountered:

Permission denied (publickey,gssapi-keyex,gssapi-with-mic,password)

Reason:

Using:

```
sudo scp ...
```

attempted root SSH login.

Rocky Linux disables direct root SSH login by default.

Resolution:

Use normal user SCP.

11.19 Copy Munge Key to Login Node

```
scp /tmp/munge.key login:/tmp/
```

11.20 Copy Munge Key to Compute01

```
scp /tmp/munge.key compute01:/tmp/
```

11.21 Copy Munge Key to Compute02

```
scp /tmp/munge.key compute02:/tmp/
```

11.22 Copy Munge Key to Compute03

```
scp /tmp/munge.key compute03:/tmp/
```

11.23 Remove Temporary Key Copy

On master:

```
sudo rm -f /tmp/munge.key
```

11.24 Install Munge Key on Client Nodes

Perform on:

- login
- compute01
- compute02
- compute03

11.25 Stop Munge Service

```
sudo systemctl stop munge
```

11.26 Install Key

```
sudo mv /tmp/munge.key /etc/munge/
```

11.27 Fix Ownership

```
sudo chown munge:munge /etc/munge/munge.key
```

11.28 Fix Permissions

```
sudo chmod 400 /etc/munge/munge.key
```

11.29 Start Munge

On ALL nodes:

```
sudo systemctl enable --now munge
```

11.30 IMPORTANT Issue Encountered — Munge Not Running on Master

Problem:

```
munge: Error: Failed to access "/var/run/munge/munge.socket.2"
```

Reason:

Munge service never started on master.

Resolution:

```
sudo systemctl enable --now munge
```

Important lesson:

Difference between:

State	Meaning
inactive	service not started
failed	service crashed

11.31 Verify Munge Service

```
systemctl status munge
```

Expected:

```
active (running)
```

11.32 Verify Munge Authentication

From master:

```
munge -n | ssh compute01 unmunge
```

11.33 Test Compute02

```
munge -n | ssh compute02 unmunge
```

11.34 Test Compute03

```
munge -n | ssh compute03 unmunge
```

11.35 Test Login Node

```
munge -n | ssh login unmunge
```

Expected:

STATUS: Success

11.36 HPC Concepts Learned

Munge introduced:

- cluster-wide authentication
- distributed trust model
- shared secret authentication
- daemon security
- inter-node trust validation

11.37 Configure Slurm Scheduler

Purpose:

Enable centralized resource scheduling.

IMPORTANT:

```
/etc/slurm/slurm.conf
```

must be IDENTICAL on ALL nodes.

11.38 IMPORTANT Issue Encountered — Default localhost Configuration

Problem:

```
sinfo
```

```
-> localhost
```

Reason:

Default example slurm.conf remained active.

Resolution:

Replace configuration with custom cluster topology.

Important HPC lesson:

Incorrect topology configuration is one of the most common Slurm admin problems.

11.39 Create Slurm Configuration

On master:

```
sudo vim /etc/slurm/slurm.conf
```

Replace ENTIRE file with:

```
ClusterName=hpccluster
```

```
SlurmctldHost=master
```

```
MpiDefault=none
```

```
ProctrackType=proctrack/cgroup
```

```
ReturnToService=2
```

```
SlurmctldPidFile=/run/slurmctld.pid
```

```
SlurmdPidFile=/run/slurmd.pid
```

```
SlurmdSpoolDir=/var/spool/slurmd
```

```
StateSaveLocation=/var/spool/slurmctld
```

```
SlurmUser=slurm
```

```
SwitchType=switch/none
```

```
TaskPlugin=task/affinity
```

```
SchedulerType=sched/backfill
```

```
SelectType=select/cons_tres
```

```
SelectTypeParameters=CR_Core
```

```
AuthType=auth/munge
```

```
NodeName=compute01 CPUs=2 State=UNKNOWN
```

```
NodeName=compute02 CPUs=2 State=UNKNOWN
```

```
NodeName=compute03 CPUs=2 State=UNKNOWN
```

```
PartitionName=debug Nodes=ALL Default=YES MaxTime=INFINITE State=UP
```


11.40 Configuration Explanation

Parameter	Purpose
ClusterName	cluster identifier
SlurmctldHost	controller node
NodeName	compute nodes
CPUs=2	available CPUs
PartitionName	scheduler queue

11.41 Create Slurm Runtime Directories

1. On Master

```
sudo mkdir -p /var/spool/slurmctld
```

```
sudo mkdir -p /var/spool/slurmd
```

11.42 Set Ownership

```
sudo chown slurm:slurm /var/spool/slurmctld
```

```
sudo chown slurm:slurm /var/spool/slurmd
```

11.43 Create slurmd Directories on Compute Nodes

1. Compute01

```
ssh compute01 sudo -S mkdir -p /var/spool/slurmd
```

2. Compute02

```
ssh compute02 sudo -S mkdir -p /var/spool/slurmd
```

3. Compute03

```
ssh compute03 sudo -S mkdir -p /var/spool/slurmd
```

11.44 IMPORTANT sudo Over SSH Issue

Problem:

```
sudo: a terminal is required to read the password
```

Resolution:

Use:

```
sudo -S
```

Important Linux administration concept:

sudo over non-interactive SSH requires stdin password handling.

11.45 Fix Ownership on Compute Nodes

1. Compute01

```
ssh compute01 sudo -S chown slurm:slurm /var/spool/slurmd
```

2. Compute02

```
ssh compute02 sudo -S chown slurm:slurm /var/spool/slurmd
```

3. Compute03

```
ssh compute03 sudo -S chown slurm:slurm /var/spool/slurmd
```

11.46 Verify Runtime Directories

```
ls -ld /var/spool/slurmctld
```

Expected:

```
slurm slurm
```

11.47 Distribute `slurm.conf`

1. Login

```
scp /etc/slurm/slurm.conf login:/tmp/
```

2. Compute01

```
scp /etc/slurm/slurm.conf compute01:/tmp/
```

3. Compute02

```
scp /etc/slurm/slurm.conf compute02:/tmp/
```

4. Compute03

```
scp /etc/slurm/slurm.conf compute03:/tmp/
```

11.48 Install Configuration on Nodes

On ALL nodes:

```
sudo mv /tmp/slurm.conf /etc/slurm/slurm.conf
```

11.49 Start Slurm Controller

ONLY on master:

```
sudo systemctl enable --now slurmctld
```

11.50 Verify Controller

```
systemctl status slurmctld
```

Expected:

active (running)

11.51 Start slurmd on Compute Nodes

On:

- compute01
- compute02
- compute03

```
sudo systemctl enable --now slurmd
```

11.52 Verify slurmd

```
systemctl status slurmd
```

Expected:

```
active (running)
```

11.53 Reconfigure Slurm

On master:

```
sudo scontrol reconfigure
```

11.54 IMPORTANT Issue Encountered — Nodes in DRAIN State

Problem:

Nodes entered:

DRAIN

after CPU topology changes.

Reason:

Slurm detected mismatch between:

- configured CPUs
- detected hardware topology

Resolution:

Clear stale slurmd state.

11.55 Stop slurmd

On compute nodes:

```
sudo systemctl stop slurmd
```

11.56 Remove Old State

```
sudo rm -rf /var/spool/slurmd/*
```

11.57 Restart slurmd

```
sudo systemctl start slurmd
```

11.58 Resume Nodes

On master:

```
sudo scontrol update NodeName=compute01 State=RESUME
```

```
sudo scontrol update NodeName=compute02 State=RESUME
```

```
sudo scontrol update NodeName=compute03 State=RESUME
```

11.59 Verify Cluster Status

```
sinfo
```

Expected:

```
PARTITION AVAIL TIMELIMIT NODES STATE NODELIST
debug* up infinite 3 idle compute[01-03]
```

11.60 Verify Node Details

```
scontrol show nodes
```

This displays:

- CPU information
- memory
- node states
- scheduler topology
- resource allocation

11.61 Test Slurm Execution

1. Single Node Test

From login node:

```
srun hostname
```

Expected:

```
compute01.local
```

11.62 Multi-node Execution

```
srun -N 3 hostname
```

Expected:

```
compute01.local
```

```
compute02.local
```

```
compute03.local
```

11.63 Multi-task Multi-node Execution

```
srun -N 3 -n 6 hostname
```

This verifies:

- task distribution
- multi-node scheduling
- CPU allocation

11.64 Multi-core Single-node Execution

```
srun -N 1 -n 2 hostname
```

Expected:

```
compute01.local
```

```
compute01.local
```

11.65 IMPORTANT Slurm Scheduling Concept

Earlier:

```
srun -N 1 -n 2 hostname
```

failed because nodes initially had:

```
CPUs=1
```

After increasing VM CPUs and updating slurm.conf:

```
CPUs=2
```

the scheduler successfully allocated 2 tasks on a single node.

Important HPC lesson:

Slurm strictly enforces:

- CPU limits
- task allocation
- topology constraints
- resource consistency

12 Shared Software Filesystem (/apps)

Goal:

Create a dedicated shared software filesystem for:

- Spack
- compilers
- MPI stacks
- scientific applications
- modulefiles
- HPC software hierarchy

This phase transitions the cluster from:

```
basic scheduler infrastructure
```

to:

```
realistic HPC software environment
```

12.1 Why Separate /apps Filesystem?

Initially:

- each VM had a 30 GB local disk
- /home was exported from master over NFS

Concern encountered:

Shared /home only had ~30 GB available.

Spack and HPC software stacks can become very large.

Important HPC storage concept:

Real HPC clusters usually separate:

Filesystem	Purpose
/home	user files
/apps	software stack
/scratch	temporary runtime data
/archive	long-term storage

Therefore a dedicated:

/apps

filesystem was created.

This is significantly closer to real HPC architecture.

12.2 IMPORTANT Storage Architecture Concept

Each VM still has its own independent local storage.

Example:

Node	Local Disk
master	30 GB
login	30 GB
compute01	30 GB
compute02	30 GB
compute03	30 GB

When:

master:/home

was mounted via NFS:

- compute node local /home became hidden
- but local disks still exist

Only the exported filesystem becomes shared.
Therefore:

/apps

was created as a dedicated shared software layer.

12.3 Add New Virtual Disk to Master

A second virtual disk was attached ONLY to:

master

Purpose:

- centralized software storage
- shared application filesystem
- Spack software stack

Disk size selected:

60 GB

Disk type:

VDI (Dynamically Allocated)

12.4 Verify New Disk Detection

After booting master:

```
lsblk
```

Expected:

```
sda    30G
sdb    60G
```

Important:

Device	Purpose
sda	OS disk
sdb	New apps disk

Only:

```
sdb
```

was modified.

12.5 Create GPT Partition Table

```
sudo fdisk /dev/sdb
```

Inside fdisk:

Create GPT label:

```
g
```

Create partition:

```
n
```

Accept defaults.

Write changes:

```
w
```

12.6 Verify Partition

```
lsblk
```

Expected:

```
sdb
sdb1
```

12.7 Create XFS Filesystem

```
sudo mkfs.xfs /dev/sdb1
```

Why XFS?

- default enterprise filesystem
- highly scalable
- excellent for large filesystems
- commonly used in HPC

12.8 Create Mount Point

```
sudo mkdir -p /apps
```

12.9 Get Filesystem UUID

```
sudo blkid /dev/sdb1
```

Example:

```
UUID="xxxxxxxx-xxxx-xxxx"
```

12.10 Configure Persistent Mount

Edit:

```
sudo vim /etc/fstab
```

Add:

```
UUID=YOUR_UUID_HERE /apps xfs defaults 0 0
```

Important Linux administration concept:

UUIDs are preferred over:

```
/dev/sdb1
```

because device names can change after reboot.

12.11 Mount Filesystem

```
sudo mount -a
```

12.12 Verify Mount

```
df -h
```

Expected:

```
/apps
```

mounted with approximately:

```
60 GB
```

12.13 Configure Ownership

```
sudo chown hpcadmin:hpcadmin /apps
```

12.14 Create Shared HPC Software Hierarchy

```
mkdir -p /apps/{spack,modulefiles,software,sources}
```

Resulting structure:

```
/apps
  spack
  modulefiles
  software
  sources
```

12.15 HPC Software Hierarchy Concept

Purpose of directories:

Directory	Purpose
spack	Spack package manager
modulefiles	custom modulefiles
software	installed software
sources	downloaded sources

This is similar to production HPC software layouts.

12.16 Export /apps via NFS

Edit exports file:

```
sudo vim /etc/exports
```

Add:

```
/apps 10.10.10.0/24(rw, sync, no_root_squash)
```

12.17 Apply NFS Exports

```
sudo exportfs -rav
```

12.18 Verify Exports

```
showmount -e localhost
```

Expected:

```
/apps
```

```
/home
```

12.19 Create Mount Points on Client Nodes

On:

- login
- compute01
- compute02

- compute03

Create:

```
sudo mkdir -p /apps
```

12.20 Configure Persistent NFS Mounts

Edit:

```
sudo vim /etc/fstab
```

Add:

```
master:/apps /apps nfs defaults,_netdev 0 0
```

12.21 Mount Shared /apps Filesystem

```
sudo mount -a
```

12.22 Verify Shared Mount

```
df -h
```

Expected:

```
master:/apps
```

mounted on all nodes.

12.23 Verify Shared Filesystem Functionality

On master:

```
touch /apps/testfile
```

On compute node:

```
ls /apps
```

Expected:

```
testfile
```

This verifies:

- shared software filesystem
- NFS functionality
- cluster-wide software visibility

12.24 HPC Architecture After /apps Setup

Current cluster storage architecture:

Filesystem	Purpose
/home	user files
/apps	software stack
local disks	local runtime storage

This resembles real HPC environments significantly more closely.

13 Spack Software Stack + MPI Runtime Validation

Goal:

Create centralized HPC software stack using:

- Spack
- shared software filesystem
- shared compilers
- shared MPI stack
- scheduler-aware runtime execution

This phase validates:

- software distribution
- multinode execution
- Slurm integration
- MPI runtime functionality

After completing this section you will be able to:

- install HPC software using Spack
- manage multiple compiler versions
- manage multiple MPI versions

- load software dynamically
- compile MPI applications
- execute distributed jobs using Slurm
- monitor and debug jobs
- understand performance metrics

13.1 IMPORTANT HPC Software Stack Concept

Instead of:

```
dnf install openmpi
```

a centralized HPC software stack was created.

Benefits:

- software installed once
- accessible cluster-wide
- compiler hierarchy support
- reproducible builds
- module-based software environments
- centralized software management

This approach is significantly closer to real HPC centers.

In production HPC environments:

- operating system packages are separated
- scientific software stack is separated
- compiler stacks are separated
- MPI stacks are separated

This separation improves:

- portability
- reproducibility
- scalability
- maintenance

13.2 IMPORTANT Shared Filesystem Concept

The cluster uses:

`/apps`

as shared application storage.

Because:

`/apps`

is mounted cluster-wide through NFS:

- software installed once becomes visible everywhere
- compute nodes automatically access installed software
- users share common software stack
- recompilation is unnecessary on compute nodes

This is one of the MOST important HPC architecture concepts.

13.3 Install Development Tools

Installed manually on ALL nodes:

```
sudo dnf groupinstall -y "Development Tools"
```

Purpose:

Provides:

- gcc
- g++
- gfortran
- make
- linker tools
- build utilities
- development environment

Important note:

System compilers are still required even when using Spack.

Reason:

Spack builds software from source code and therefore still requires:

- assembler
- linker
- basic compiler toolchain

13.4 Move into Shared Application Filesystem

Run on:

```
master
```

Move into shared applications directory:

```
cd /apps
```

Verify:

```
pwd
```

Expected:

```
/apps
```

13.5 Clone Spack into Shared Filesystem

Clone Spack:

```
git clone https://github.com/spack/spack.git
```

Expected:

```
/apps/spack
```

This creates centralized package manager installation.

13.6 Enter Spack Directory

```
cd spack
```

Verify:

```
pwd
```

Expected:

```
/apps/spack
```

13.7 Initialize Spack Environment

Initialize Spack:

```
source share/spack/setup-env.sh
```

Purpose:

Adds:

- spack command
- shell integration
- environment variables
- package management environment

to current shell session.

Without initialization:

- spack command will not work
- spack load will fail
- software environments will not activate

13.8 IMPORTANT Shell Environment Concept

`source`

does NOT execute script normally.
Instead:

it modifies current shell environment

This is extremely important in:

- Spack
- modules
- MPI environments
- compiler environments

13.9 Make Spack Permanent

To avoid sourcing manually every session:
Edit:

`vim ~/.bashrc`

Add:

`source /apps/spack/share/spack/setup-env.sh`

Reload shell:

`source ~/.bashrc`

Verify:

`spack --version`

Expected:

- Spack version information

13.10 Clone Additional Spack Repository

Clone repository:

```
git clone https://github.com/spack/spack-packages.git
```

Purpose:

Adds additional package definitions.

13.11 Configure Spack Repository

```
spack repo set --destination "$(pwd)/spack-packages" builtin
```

Purpose:

Configure additional repository support.

This allows:

- package metadata
- package recipes
- software definitions

to become visible to Spack.

13.12 Verify Spack Functionality

```
spack info gcc
```

Expected:

Detailed GCC package information.

This verifies:

- Spack environment working
- package repository functional
- package metadata accessible

13.13 IMPORTANT HPC Software Stack Concept

Spack operates as:

source-based HPC package manager

Unlike:

Package Manager	Purpose
dnf	operating system
Spack	HPC software ecosystem

Production HPC systems commonly separate:

- OS packages
- scientific software stacks

exactly like this setup.

13.14 IMPORTANT Spack Build Concept

When installing software using Spack:

software is compiled from source

Advantages:

- architecture optimization
- compiler flexibility
- dependency control
- reproducibility
- variant support

Disadvantages:

- longer installation time

13.15 Detect Available Compilers

Before installing software:

```
spack compiler find
```

Purpose:

Detect system compilers available for building software.

Verify detected compilers:

```
spack compilers
```

Expected:

- compiler list

13.16 Install GCC Using Spack

Install GCC version:

```
spack install gcc@13.4.0
```

Purpose:

Create dedicated HPC compiler stack.

This may take considerable time because:

- GCC itself is compiled from source
- dependencies are built
- optimization occurs

13.17 IMPORTANT Multiple Compiler Version Concept

Production HPC systems frequently maintain multiple compiler versions simultaneously.

Example:

```
gcc@11  
gcc@12  
gcc@13  
intel  
nvhpc  
clang
```

Reasons:

- application compatibility
- compiler ABI differences
- reproducibility
- benchmark consistency

13.18 Verify Installed Software

List installed packages:

```
spack find
```

Expected example:

```
==> 1 installed package  
gcc@13.4.0
```

13.19 Load Installed GCC Compiler

```
spack load gcc@13.4.0
```

Purpose:

Activate compiler environment.

Verify compiler path:

```
which gcc
```

Expected:

- Spack GCC path

Verify version:

```
gcc --version
```

Expected:

- GCC 13.4.0

13.20 IMPORTANT Runtime Environment Concept

When software is loaded:

- PATH changes
- LD_LIBRARY_PATH changes
- MANPATH changes

This allows applications to locate:

- executables
- shared libraries
- runtime dependencies

automatically.

13.21 Register Newly Installed Compiler

After installing GCC:

```
spack compiler find
```

Purpose:

Allow Spack to use newly installed GCC for future package builds.

Verify:

```
spack compilers
```

Expected:

- new GCC compiler entry

13.22 Install OpenMPI Using Spack

Install MPI runtime:

```
spack install openmpi
```

Purpose:

Create shared MPI runtime stack.

This installation includes:

- MPI runtime
- MPI compiler wrappers
- MPI libraries
- distributed runtime environment

13.23 IMPORTANT MPI Compiler Dependency Concept

MPI stacks are tightly coupled with compiler toolchains.

Example:

`openmpi` built using `gcc@13`

This is important because:

- ABI compatibility matters
- runtime compatibility matters
- mixed compiler environments may fail

13.24 Verify Installed Software Tree

Show dependency tree:

`spack find -d`

Purpose:

Displays:

- dependency hierarchy
- compiler hierarchy
- runtime dependencies

Example:

```
openmpi
  ^gcc@13.4.0
```

13.25 IMPORTANT Spack Spec Concept

Every Spack installation contains:

- package version
- compiler information
- dependency tree
- unique installation hash

Example:

`openmpi@5.0.7%gcc@13.4.0`

Meaning:

Component	Meaning
openmpi	package
5.0.7	package version
gcc@13.4.0	compiler used

13.26 IMPORTANT Hash-Based Software Loading

Sometimes multiple builds of same version exist because of:

- different compilers
- different variants
- different dependencies

Spack assigns:

unique hashes

View hashes:

`spack find -l`

Expected:

`abc1234 gcc@13.4.0`

`xyz5678 openmpi@5.0.7`

Load using hash:

`spack load /xyz5678`

This is extremely common in production HPC systems.

13.27 Load OpenMPI Environment

```
spack load openmpi
```

Verify:

```
which mpicc
```

Expected:

- Spack OpenMPI wrapper path

Verify MPI runtime:

```
mpirun --version
```

Expected:

- OpenMPI version information

13.28 IMPORTANT MPI Wrapper Concept

```
mpicc
```

is NOT a compiler itself.

It is a wrapper that automatically injects:

- MPI headers
- MPI libraries
- runtime linker flags

around GCC.

This is a critical HPC build concept.

13.29 Verify Loaded Software

Display loaded packages:

```
spack find --loaded
```

Unload OpenMPI:

```
spack unload openmpi
```

Unload all software:

```
spack unload --all
```

Verify:

```
spack find --loaded
```

Expected:

- no loaded packages

13.30 Create MPI Performance Validation Program

Create source file:

```
vim mpi_sum.c
```

Purpose:

- validate MPI runtime
- validate multinode execution
- validate scheduler-aware execution
- compare serial vs parallel runtime
- demonstrate reduction operations
- calculate speedup
- calculate efficiency

Paste:

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>

#define N 1000000000

double serial_sum()
{
```

```

    double sum = 0.0;

    for(long long int i = 0; i < N; i++)
    {
        sum += i;
    }

    return sum;
}

int main(int argc, char *argv[])
{
    int rank, size;

    long long int start_index;
    long long int end_index;
    long long int chunk;

    double local_sum = 0.0;
    double global_sum = 0.0;

    double serial_start;
    double serial_end;

    double parallel_start;
    double parallel_end;

    MPI_Init(&argc, &argv);

    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    if(rank == 0)
    {
        serial_start = MPI_Wtime();

        double serial_result = serial_sum();

        serial_end = MPI_Wtime();
    }
}

```

```

        printf("\n");
        printf("===== SERIAL COMPUTATION =====\n");
        printf("Sum = %.0f\n", serial_result);
        printf("Time taken = %f seconds\n",
                serial_end - serial_start);
    }

    MPI_Barrier(MPI_COMM_WORLD);

    chunk = N / size;

    start_index = rank * chunk;

    if(rank == size - 1)
        end_index = N;
    else
        end_index = start_index + chunk;

    parallel_start = MPI_Wtime();

    for(long long int i = start_index;
        i < end_index;
        i++)
    {
        local_sum += i;
    }

    MPI_Reduce(&local_sum,
               &global_sum,
               1,
               MPI_DOUBLE,
               MPI_SUM,
               0,
               MPI_COMM_WORLD);

    parallel_end = MPI_Wtime();

    if(rank == 0)
    {
        double parallel_time =

```

```

        parallel_end - parallel_start;

double serial_time =
    serial_end - serial_start;

printf("\n");
printf("===== PARALLEL COMPUTATION =====\n");
printf("Processes used = %d\n", size);
printf("Sum = %.0f\n", global_sum);
printf("Time taken = %f seconds\n",
        parallel_time);

printf("\n");
printf("===== PERFORMANCE =====\n");
printf("Speedup = %f\n",
        serial_time / parallel_time);

printf("Efficiency = %f\n",
        (serial_time / parallel_time) / size);
}

MPI_Finalize();

return 0;
}

```

13.31 IMPORTANT MPI Reduction Concept

This program uses:

`MPI_Reduce()`

Purpose:

Combine partial sums from all MPI ranks into:

single global result

Flow:

local sums

↓

`MPI_Reduce`

↓

`global sum`

This is one of the MOST important MPI collective communication operations.

13.32 IMPORTANT Parallel Execution Concept

Serial execution:

`single process computes entire workload`

Parallel execution:

`multiple MPI ranks divide workload`

Example:

Rank	Work
rank0	0 -> 25%
rank1	25% -> 50%
rank2	50% -> 75%
rank3	75% -> 100%

Advantages:

- reduced execution time
- distributed workload
- scalable computation

13.33 Load Software Before Compilation

Initialize Spack:

`source /apps/spack/share/spack/setup-env.sh`

Load compiler:

`spack load gcc@13.4.0`

Load MPI runtime:

```
spack load openmpi
```

Verify:

```
which mpicc
```

Expected:

- MPI compiler wrapper path

13.34 Compile MPI Program

Compile:

```
mpicc mpi_sum.c -o mpi_sum
```

Verify executable:

```
ls
```

Expected:

```
mpi_sum
```

13.35 IMPORTANT MPI Compilation Concept

MPI programs are compiled using:

```
mpicc
```

instead of:

```
gcc
```

because MPI wrapper automatically injects:

- MPI headers
- MPI libraries
- linker flags

13.36 Interactive MPI Runtime Validation

Run interactively:

```
mpirun -np 4 ./mpi_sum
```

Expected:

```
===== SERIAL COMPUTATION =====  
...  
  
===== PARALLEL COMPUTATION =====  
...  
  
===== PERFORMANCE =====  
...
```

This validates:

- MPI runtime
- MPI communication
- reduction operations
- distributed execution

13.37 IMPORTANT Shared Filesystem Advantage

Because:

- source code
- executable
- software stack

exist on shared filesystem:

```
/apps  
or  
/home
```

compute nodes automatically access:

- compiled binaries
- runtime libraries
- MPI stack

without recompilation.

This is one of the major HPC architecture advantages.

13.38 Create Slurm Batch Script

Create batch script:

```
vim submit.sh
```

Paste:

```
#!/bin/bash
#SBATCH -N 3
#SBATCH --ntasks-per-node=1
#SBATCH --job-name=test
#SBATCH --output=%J.out
#SBATCH --error=%J.err
#SBATCH --time=01:00:00
##SBATCH --odelist=compute03,compute01

source /apps/spack/share/spack/setup-env.sh

spack load openmpi

mpirun -np $SLURM_NTASKS ./$1

Make executable:

chmod +x submit.sh
```

13.39 IMPORTANT Slurm Script Explanation

Line	Purpose
#SBATCH -N 3	request 3 nodes
--ntasks-per-node=1	one MPI rank per node
--job-name=test	job name
--output=%J.out	stdout file
--error=%J.err	stderr file
--time=01:00:00	walltime limit

13.40 IMPORTANT Runtime Environment Section

Inside batch script:

```
source /apps/spack/share/spack/setup-env.sh
```

initializes Spack environment.

```
spack load openmpi
```

loads MPI runtime stack.

Without these:

- mpirun may fail
- MPI libraries may not load
- runtime linker errors may occur

13.41 IMPORTANT Slurm Runtime Variable

`$SLURM_NTASKS`

contains:

total MPI process count allocated by Slurm

Example:

Nodes	Tasks per node	Total
3	1	3

Therefore:

```
mpirun -np $SLURM_NTASKS
```

automatically launches correct number of MPI ranks.

13.42 Submit MPI Job

Submit job:

```
sbatch submit.sh mpi_sum
```

Expected:

```
Submitted batch job 101
```

This validates:

- Slurm scheduler
- distributed allocation
- multinode MPI execution
- scheduler-aware runtime launch

13.43 Monitor Jobs

Check user jobs:

```
squeue --me
```

Expected:

```
JOBID PARTITION NAME USER ST TIME NODES NODELIST(REASON)
```

13.44 IMPORTANT Job States

State	Meaning
PD	pending
R	running
CG	completing
CD	completed
F	failed

13.45 View Detailed Job Information

Show detailed job info:

```
scontrol show job <jobid>
```

Example:

```
scontrol show job 101
```

Displays:

- allocated nodes
- CPUs
- runtime
- scheduling information
- execution directory
- output locations

13.46 Check Cluster Status

Check cluster nodes:

```
sinfo
```

Expected:

```
debug* up infinite 3 idle compute[01-03]
```

13.47 IMPORTANT Node States

State	Meaning
idle	available
alloc	allocated
mix	partially used
drain	unavailable
down	offline

13.48 Check Output Files

After completion:

```
ls
```

Expected:

```
101.out
```

```
101.err
```

View output:

```
cat 101.out
```

Expected:

- serial runtime
- parallel runtime
- speedup
- efficiency

13.49 IMPORTANT HPC Performance Metrics

The program calculates:

Metric	Meaning
Speedup	$\text{serial_time} / \text{parallel_time}$
Efficiency	$\text{speedup} / \text{processes}$

Ideal values:

Metric	Ideal
Speedup	near process count
Efficiency	near 1.0

Real systems experience:

- communication overhead
- synchronization cost
- MPI latency
- reduction overhead

13.50 Cancel Jobs

Cancel job:

```
scancel <jobid>
```

Example:

```
scancel 101
```

Cancel all user jobs:

```
scancel -u $USER
```

13.51 Interactive Slurm Execution

Launch interactive shell:

```
srun --pty bash
```

Distributed hostname test:

```
srun -N 3 -n 3 hostname
```

Expected:

- hostnames from allocated nodes

This validates:

- scheduler allocation
- distributed execution
- runtime communication

13.52 IMPORTANT HPC Workflow

Typical HPC workflow:

```
Load software
  ↓
Compile application
  ↓
Prepare Slurm script
  ↓
Submit job
  ↓
Monitor execution
  ↓
Analyze outputs
```

13.53 IMPORTANT Production HPC Concept

Production HPC systems usually separate:

- operating system packages
- scientific software stack
- compiler stack
- MPI runtime stack
- scheduler integration

exactly like this setup.

This architecture provides:

- reproducibility
- scalability
- centralized administration
- software consistency
- runtime portability

14 Centralized Identity Management using OpenLDAP + nslcd

Goal:

Transform the cluster from:

`local-user-only infrastructure`

into:

`centralized multi-user HPC environment`

using:

- OpenLDAP
- nslcd
- NSS
- PAM

This phase introduces centralized user management for the HPC cluster.
After completion:

- users are created once centrally
- login node recognizes LDAP users
- compute nodes recognize LDAP users
- centralized authentication works cluster-wide
- automatic home directory creation works
- Slurm can later use centralized identities

14.1 Important Architecture

Cluster roles:

Node	Role
master	LDAP server
login	LDAP client
compute01	LDAP client
compute02	LDAP client
compute03	LDAP client

Important concept:

Only master runs LDAP SERVER (slapd).

Login and compute nodes run LDAP CLIENTS (nslcd).

14.2 Important Linux Authentication Concepts

Before starting, understand the layers.

Component	Purpose
LDAP	centralized user database
NSS	user/group lookup
PAM	authentication
nslcd	bridge between Linux and LDAP

Authentication flow:

```
SSH Login
↓
PAM
↓
nslcd
↓
LDAP server
```

14.3 Install OpenLDAP Server on Master Node

Run ONLY on:

master

Install LDAP packages:

```
sudo dnf install -y \  
openldap \  
openldap-servers \  
openldap-clients
```

Package explanation:

Package	Purpose
openldap	LDAP utilities
openldap-servers	slapd server daemon
openldap-clients	ldapsearch, ldapadd etc

14.4 Enable LDAP Server

Run ONLY on:

master

Enable and start slapd:

```
sudo systemctl enable --now slapd
```

Verify:

```
systemctl status slapd
```

Expected:

active (running)

14.5 Generate LDAP Admin Password Hash

Run ONLY on:

master

Generate secure LDAP admin password hash:

```
slappasswd
```

Enter desired password.

Example:

cluster123

Example output:

{SSHA}xxxxxxxxxxxxxxxxxx

Important:

Copy this hash carefully.

It will be used in LDAP database configuration.

14.6 Configure LDAP Database

Run ONLY on:

master

Create LDAP database configuration file:

vim ~/db.ldif

Paste:

```
dn: olcDatabase={2}mdb,cn=config
changetype: modify
replace: olcSuffix
olcSuffix: dc=hpc,dc=local
```

```
dn: olcDatabase={2}mdb,cn=config
changetype: modify
replace: olcRootDN
olcRootDN: cn=Manager,dc=hpc,dc=local
```

```
dn: olcDatabase={2}mdb,cn=config
changetype: modify
replace: olcRootPW
olcRootPW: PASTE_HASH_HERE
```

Replace:

PASTE_HASH_HERE

with the hash generated using:

slappasswd

14.7 Important LDAP Concepts

LDAP uses a hierarchical tree called:

DIT -- Directory Information Tree

Example structure:

```
dc=hpc,dc=local
ou=People
ou=Groups
```

Meaning of terms:

Term	Meaning
dc	domain component
ou	organizational unit
cn	common name

14.8 Apply LDAP Database Configuration

Run ONLY on:

master

Apply configuration:

```
sudo ldapmodify -Y EXTERNAL -H ldapi:/// \
-f ~/db.ldif
```

Expected:

modifying entry

14.9 Import LDAP Schemas

Schemas define:

- user attributes
- group attributes
- UNIX account structure

Run ONLY on:

master

Import cosine schema:

```
sudo ldapadd -Y EXTERNAL -H ldapi:/// \  
-f /etc/openldap/schema/cosine.ldif
```

Import nis schema:

```
sudo ldapadd -Y EXTERNAL -H ldapi:/// \  
-f /etc/openldap/schema/nis.ldif
```

Import inetorgperson schema:

```
sudo ldapadd -Y EXTERNAL -H ldapi:/// \  
-f /etc/openldap/schema/inetorgperson.ldif
```

Schema explanation:

Schema	Purpose
cosine	basic directory attributes
nis	UNIX/Linux account support
inetorgperson	user/person identity support

14.10 Create LDAP Base Tree

Run ONLY on:

master

Create base LDAP tree:


```
vim ~/base.ldif
```

Paste:

```
dn: dc=hpc,dc=local
objectClass: top
objectClass: dcObject
objectClass: organization
o: HPC Cluster
dc: hpc

dn: ou=People,dc=hpc,dc=local
objectClass: organizationalUnit
ou: People

dn: ou=Groups,dc=hpc,dc=local
objectClass: organizationalUnit
ou: Groups
```

Apply tree:

```
ldapadd -x \
-D "cn=Manager,dc=hpc,dc=local" \
-W \
-f ~/base.ldif
```

Enter LDAP admin password.

14.11 Verify LDAP Tree

Run ONLY on:

master

Verify tree:

```
ldapsearch -x -b "dc=hpc,dc=local"
```

Expected:

- dc=hpc,dc=local
- ou=People
- ou=Groups

14.12 Create LDAP Group

Run ONLY on:

master

Create group file:

```
vim ~/group.ldif
```

Paste:

```
dn: cn=hpcusers,ou=Groups,dc=hpc,dc=local
objectClass: posixGroup
cn: hpcusers
gidNumber: 10000
```

Add group:

```
ldapadd -x \
-D "cn=Manager,dc=hpc,dc=local" \
-W \
-f ~/group.ldif
```

14.13 Create LDAP User Password Hash

Run ONLY on:

master

Generate password hash:

```
slappasswd
```

Example password:

```
test123
```

Copy generated hash carefully.

14.14 Create LDAP User

Run ONLY on:

master

Create user file:

```
vim ~/user.ldif
```

Paste:

```
dn: uid=testUser,ou=People,dc=hpc,dc=local
objectClass: inetOrgPerson
objectClass: posixAccount
objectClass: shadowAccount
cn: Test User
sn: User
uid: testUser
uidNumber: 10001
gidNumber: 10000
homeDirectory: /home/testUser
loginShell: /bin/bash
userPassword: PASTE_HASH_HERE
```

Replace:

PASTE_HASH_HERE

with actual hash from:

```
slappasswd
```

Add user:

```
ldapadd -x \
-D "cn=Manager,dc=hpc,dc=local" \
-W \
-f ~/user.ldif
```

14.15 Important LDAP User Attributes

Attribute	Purpose
uid	username
uidNumber	UNIX UID
gidNumber	primary group
homeDirectory	user home
loginShell	login shell
cn	common name
sn	surname

14.16 Verify LDAP User

Run ONLY on:

master

Verify user:

```
ldapsearch -x \  
-b "dc=hpc,dc=local" uid=testUser
```

Expected:

- LDAP user entry

14.17 Configure LDAP Client on Login Node

Now begin client-side integration.

Run ONLY on:

login

Install LDAP client packages:

```
sudo dnf install -y \  
openldap-clients \  
nss-pam-ldapd
```

14.18 Configure nslcd

Run ONLY on:

login

Create configuration:

```
sudo vim /etc/nslcd.conf
```

Paste:

```
uid nslcd
gid ldap
```

```
uri ldap://master
base dc=hpc,dc=local
```

Explanation:

Parameter	Purpose
uri	LDAP server
base	LDAP search root

Secure configuration:

```
sudo chmod 600 /etc/nslcd.conf
```

14.19 Enable nslcd

Run ONLY on:

login

Enable service:

```
sudo systemctl enable --now nslcd
```

Verify:

```
systemctl status nslcd
```

Expected:

```
active (running)
```

14.20 Verify LDAP Connectivity

Run ONLY on:

login

Test LDAP communication:

```
ldapsearch -x \  
-H ldap://master \  
-b "dc=hpc,dc=local" uid=testUser
```

Expected:

- LDAP user entry

This verifies:

- login node can communicate with master LDAP server

14.21 Configure NSS

Run ONLY on:

login

Edit:

```
sudo vim /etc/nsswitch.conf
```

Modify lines:

```
passwd: files ldap  
group: files ldap  
shadow: files ldap
```

Important:

This tells Linux:

check local files first, then LDAP.

14.22 Restart nslcd

Run ONLY on:

login

Restart service:

```
sudo systemctl restart nslcd
```

14.23 Verify LDAP User Resolution

Run ONLY on:

login

Test centralized user lookup:

```
getent passwd testUser
```

Expected:

```
testUser:x:10001:10000:Test User:/home/testUser:/bin/bash
```

Important concept:

This proves Linux itself recognizes LDAP users.

This is one of the most important LDAP verification steps.

14.24 NOTE (Do this for master too)

Now we must configure: **master** with EXACT same setup (master should be able to recognise the user).

14.25 Install Automatic Home Directory Support

Run ONLY on:

login

Install package:

```
sudo dnf install -y oddjob-mkhomedir
```

Enable service:

```
sudo systemctl enable --now oddjobd
```

Purpose:

Automatically creates user home directory during first successful login.

14.26 Configure PAM Authentication

Run ONLY on:

login

Edit:

```
sudo vim /etc/pam.d/system-auth
```

Add LDAP lines in appropriate sections.

Add in auth section:

```
auth          sufficient    pam_ldap.so use_first_pass
```

Add in account section:

```
account       sufficient    pam_ldap.so
```

Add in password section:

```
password      sufficient    pam_ldap.so use_authtok
```

Add in session section:

```
session       optional      pam_ldap.so
session       optional      pam_mkhomedir.so
```

Important:

Do NOT remove pam_unix.so lines.

14.27 Configure Second PAM File

Run ONLY on:

login

Edit:

```
sudo vim /etc/pam.d/password-auth
```

Add SAME LDAP PAM lines.

Reason:

Different login methods use different PAM stacks.

14.28 Restart Services

Run ONLY on:

login

Restart services:

```
sudo systemctl restart nslcd
sudo systemctl restart oddjobd
```

14.29 Verify LDAP Authentication

Run ONLY on:

login

Test login:

```
su - testUser
```

Enter LDAP password.

Expected:

Creating directory '/home/testUser'

Successful login indicates:

- LDAP authentication working
- PAM integration working
- automatic home directory creation working

14.30 Verify User Home

Run ONLY on:

```
login
```

After login:

```
pwd
```

Expected:

```
/home/testUser
```

1. NOW WE MUST REPLICATE TO COMPUTE NODES

Currently:

```
login node works fully
```

Now we configure:

- compute01
- compute02
- compute03

with EXACT same client setup.

14.31 Troubleshooting

1. Important Slurm + LDAP Troubleshooting

This section explains one of the MOST common HPC setup issues.

2. Symptom

LDAP users:

- can SSH login
- can use su
- appear in getent passwd

BUT:

```
sbatch submit.sh
```

results in:

- jobs stuck in pending
- BeginTime state
- scheduler instability
- sinfo hanging
- Slurm credential errors
- slurmctld crashes

3. Root Cause

Slurm internally works using:

- UID
- GID
- NSS identity resolution

When launching jobs, Slurm performs internal user identity lookups.

If master node cannot resolve LDAP users correctly through:

- NSS
- nslcd
- LDAP

then:

- Slurm credential generation fails
- scheduler may crash
- jobs cannot launch

Common error:

```
_fill_cred_gids: getpwuid failed for uid=10001
```

Meaning:

Slurm cannot resolve LDAP user identity correctly.

4. Important Verification Commands

Run on:

```
master  
login  
compute nodes
```

Verify LDAP resolution:

```
getent passwd testUser  
id testUser
```

Expected:

```
uid=10001(testUser) gid=10000(hpcusers)
```

These commands MUST work correctly on ALL cluster nodes.

5. Important Master Node Verification

The master node MUST also resolve LDAP users correctly.

Verify:

```
getent passwd testUser
```

If master cannot resolve LDAP users:

- slurmctld may fail
- scheduler may crash
- jobs remain pending

Verify nslcd configuration on master:

```
sudo vim /etc/nslcd.conf
```

Example:

```
uid nslcd
```

```
gid ldap
```

```
uri ldap://master
```

```
base dc=hpc,dc=local
```

Verify NSS configuration:

```
sudo vim /etc/nsswitch.conf
```

Ensure:

```
passwd: files ldap
```

```
group: files ldap
```

```
shadow: files ldap
```

Restart LDAP client:

```
sudo systemctl restart nslcd
```

6. Important Slurm Architecture

Correct architecture:

Node	Service
master	slurmctld
login	login services only
compute01	slurmd
compute02	slurmd
compute03	slurmd

IMPORTANT:

Login node should NOT run slurmd.

Because login node is NOT configured as a compute node in:

```
/etc/slurm/slurm.conf
```

Attempting to start slurmd on login gives:

```
slurmd: fatal: Unable to determine this slurmd's NodeName
```

This is NORMAL and expected.

Correct action on login node:

```
sudo systemctl stop slurmd
sudo systemctl disable slurmd
```

Expected afterward:

```
inactive (dead)
```

This is CORRECT for login nodes.

7. Important Slurm Recovery Commands

Whenever:

- LDAP users are added
- NSS configuration changes
- PAM configuration changes
- nslcd configuration changes

restart Slurm services.

Restart controller on master:

```
sudo systemctl restart slurmctld
```

Restart compute daemons on compute nodes:

```
sudo systemctl restart slurmd
```

Restart LDAP client:

```
sudo systemctl restart nslcd
```

Verify cluster:

```
sinfo
```

Expected:

```
debug* up infinite 3 idle compute[01-03]
```

8. Important Scheduler Debugging

If:

```
sinfo
```

hangs indefinitely:

Possible cause:

```
slurmctld crashed
```

Verify:

```
systemctl status slurmctld
```

or:

```
journalctl -u slurmctld -n 50
```

14.32 Current Cluster State After LDAP Setup

Cluster now contains:

Component	Status
OpenLDAP server	working
centralized users	working
centralized authentication	working
nsld integration	working
NSS integration	working
PAM authentication	working
automatic home creation	working

14.33 Important Debugging Commands

Verify LDAP tree:

```
ldapsearch -x -b "dc=hpc,dc=local"
```

Verify LDAP user:

```
ldapsearch -x -b "dc=hpc,dc=local" uid=testUser
```

Verify LDAP authentication:

```
ldapwhoami -x \  
-D "uid=testUser,ou=People,dc=hpc,dc=local" \  
-W \  
-H ldap://master
```

Verify Linux user resolution:

```
getent passwd testUser
```

Verify nslcd service:

```
systemctl status nslcd
```

14.34 Important Architecture Understanding

Master node:

- provides centralized identity database

Login/compute nodes:

- authenticate users
- resolve LDAP users
- create user sessions

This separation is fundamental in distributed authentication systems.

15 Automated LDAP User Creation Script

Managing users manually using individual LDIF files becomes difficult as the number of users increases.

To simplify user management, an automation script can be used to:

- create LDAP group
- create LDAP user
- generate password automatically
- generate password hash automatically
- add entries into LDAP

This script automates complete LDAP user creation.

IMPORTANT:

The script creates:

- private group for user
- LDAP user entry
- hashed LDAP password

automatically.

15.1 Script Location

Example:

```
cd ~/ldap-scripts  
ls
```

Expected:

```
createUser.sh
```

15.2 User Creation Script

Create file:

```
vim createUser.sh
```

Paste the complete script exactly as shown below:

```
#!/bin/bash

LDAP_ADMIN="cn=Manager,dc=hpc,dc=local"
BASE_DN="dc=hpc,dc=local"

USERNAME=$1
UIDNUM=$2
GIDNUM=$3

if [ $# -ne 3 ]; then
    echo "Usage: $0 <username> <uid> <gid>"
    exit 1
fi

# Check if user already exists
USER_EXISTS=$(ldapsearch -x \
-b "ou=People,$BASE_DN" "(uid=${USERNAME})" dn \
| grep "^dn:")

if [ ! -z "$USER_EXISTS" ]; then
    echo "ERROR: User ${USERNAME} already exists!"
    exit 1
fi

# Check if group already exists
GROUP_EXISTS=$(ldapsearch -x \
-b "ou=Groups,$BASE_DN" "(cn=${USERNAME})" dn \
| grep "^dn:")

if [ ! -z "$GROUP_EXISTS" ]; then
    echo "ERROR: Group ${USERNAME} already exists!"
    exit 1
fi
```

```

PASSWORD="${USERNAME}@0123"

PASSWORD_HASH=$(slappasswd -s "$PASSWORD")

echo
echo "=====
echo "Creating LDAP User"
echo "=====
echo "Username : ${USERNAME}"
echo "UID      : ${UIDNUM}"
echo "GID      : ${GIDNUM}"
echo "=====

# Create private group LDIF
cat > /tmp/${USERNAME}_group.ldif <<EOF
dn: cn=${USERNAME},ou=Groups,${BASE_DN}
objectClass: posixGroup
cn: ${USERNAME}
gidNumber: ${GIDNUM}
EOF

# Create user LDIF
cat > /tmp/${USERNAME}.ldif <<EOF
dn: uid=${USERNAME},ou=People,${BASE_DN}
objectClass: inetOrgPerson
objectClass: posixAccount
objectClass: shadowAccount
cn: ${USERNAME}
sn: ${USERNAME}
uid: ${USERNAME}
uidNumber: ${UIDNUM}
gidNumber: ${GIDNUM}
homeDirectory: /home/${USERNAME}
loginShell: /bin/bash
userPassword: ${PASSWORD_HASH}
EOF

echo
echo "Adding LDAP group..."

```

```

ldapadd -x \
-D "${LDAP_ADMIN}" \
-W \
-f /tmp/${USERNAME}_group.ldif

if [ $? -ne 0 ]; then
    echo "ERROR: Failed to create LDAP group!"
    exit 1
fi

echo
echo "Adding LDAP user..."

ldapadd -x \
-D "${LDAP_ADMIN}" \
-W \
-f /tmp/${USERNAME}.ldif

if [ $? -ne 0 ]; then
    echo "ERROR: Failed to create LDAP user!"
    exit 1
fi

echo
echo "=====
echo "LDAP User Created Successfully"
echo "=====
echo "Username : ${USERNAME}"
echo "Password : ${PASSWORD}"
echo "UID      : ${UIDNUM}"
echo "GID      : ${GIDNUM}"
echo "=====

```

15.3 Make Script Executable

Before using the script:

```
chmod +x createUser.sh
```

15.4 Understanding Script Arguments

The script requires THREE arguments:

Argument	Meaning
username	LDAP username
uid	UNIX UID
gid	UNIX GID

Syntax:

```
./createUser.sh <username> <uid> <gid>
```

Example:

```
./createUser.sh student01 11001 11001
```

Explanation:

Value	Meaning
student01	username
11001	UID
11001	GID

IMPORTANT:

The script creates a private group automatically.

Therefore username and group name become identical.

Example:

Entity	Created Automatically
User	student01
Group	student01
Home Directory	/home/student01

15.5 Default Password Format

The script automatically creates password using:

`username@@123`

Example:

Username	Password
student01	student01@@123
student02	student02@@123

IMPORTANT:

Users should change password after first login.

15.6 What the Script Automatically Does

The script performs following operations automatically:

Operation	Description
Check existing user	avoids duplicate users
Check existing group	avoids duplicate groups
Generate password hash	secure LDAP password
Create LDAP group LDIF	automatic
Create LDAP user LDIF	automatic
Add LDAP group	automatic
Add LDAP user	automatic

15.7 Example User Creation

Example:

```
./createUser.sh student02 11002 11002
```

Expected output:

```
=====
Creating LDAP User
=====
Username : student02
UID      : 11002
GID      : 11002
=====
```

LDAP admin password will be requested.

After successful creation:

```
=====
LDAP User Created Successfully
=====
Username : student02
Password : student02@@123
UID      : 11002
GID      : 11002
=====
```

15.8 Verify User Creation

Verify user using:

```
getent passwd student02
```

or:

```
id student02
```

Expected:

```
uid=11002(student02)
```

Verify LDAP entry:

```
ldapsearch -x \
-b "dc=hpc,dc=local" uid=student02
```

15.9 Verify Login

Test login:

```
su - student02
```

or:

```
ssh student02@login
```

Expected:

```
successful login
```

15.10 Important UID/GID Notes

IMPORTANT:

UID and GID values must remain unique cluster-wide.

Avoid:

- duplicate UID
- duplicate GID

because duplicate IDs can create:

- ownership conflicts
- wrong usernames
- incorrect file ownership
- Slurm identity problems

Example of BAD configuration:

User	UID
student01	11001
student02	11001

This can cause:

`whoami` returning wrong username

Always use unique UID/GID values.

15.11 Important HPC Identity Concept

In HPC clusters:

- NFS
- Slurm
- MPI
- LDAP

- SSH

all depend on:

`consistent unique UID/GID mappings`

across the entire cluster.

This is one of the MOST important concepts in centralized HPC authentication.

16 Automated LDAP User Deletion Script

Managing users manually using:

- `ldapdelete`
- group deletion
- home directory cleanup

can become repetitive and error-prone.

To simplify cleanup operations, an automated LDAP deletion script can be used.

This script automates:

- LDAP user deletion
- optional LDAP group deletion
- optional home directory deletion

IMPORTANT:

The script provides interactive cleanup options for safer user management.

16.1 Script Location

Example:

```
cd ~/ldap-scripts
ls
```

Expected:

```
deleteUser.sh
```

16.2 User Deletion Script

Create file:

```
vim deleteUser.sh
```

Paste the complete script exactly as shown below:

```
#!/bin/bash

LDAP_ADMIN="cn=Manager,dc=hpc,dc=local"
BASE_DN="dc=hpc,dc=local"

USERNAME=$1

if [ $# -ne 1 ]; then
    echo "Usage: $0 <username>"
    exit 1
fi

echo
echo "=====
echo "Deleting LDAP User"
echo "=====
echo "Username : ${USERNAME}"
echo "=====

echo
echo "Deleting LDAP user entry..."

ldapdelete -x \
-D "${LDAP_ADMIN}" \
-W \
"uid=${USERNAME},ou=People,${BASE_DN}"

if [ $? -ne 0 ]; then
    echo "ERROR: Failed to delete LDAP user!"
    exit 1
fi

echo
```

```

echo "LDAP user deleted successfully."

# Ask about deleting group
echo
read -p "Delete LDAP group '${USERNAME}'? (y/n): " DELETE_GROUP

case $DELETE_GROUP in
    y|Y)

        ldapdelete -x \
        -D "${LDAP_ADMIN}" \
        -W \
        "cn=${USERNAME},ou=Groups,${BASE_DN}"

        if [ $? -eq 0 ]; then
            echo "LDAP group deleted successfully."
        else
            echo "WARNING: Failed to delete LDAP group."
        fi
        ;;

    n|N)
        echo "Skipping LDAP group deletion."
        ;;

    *)
        echo "Invalid option. Skipping group deletion."
        ;;
esac

# Ask about deleting home directory
echo
read -p "Delete home directory '/home/${USERNAME}'? (y/n): " DELETE_HOME

case $DELETE_HOME in
    y|Y)

        if [ -d "/home/${USERNAME}" ]; then
            sudo rm -rf "/home/${USERNAME}"

```

```

        if [ $? -eq 0 ]; then
            echo "Home directory deleted successfully."
        else
            echo "WARNING: Failed to delete home directory."
        fi
    else
        echo "Home directory does not exist."
    fi
;;

n|N)
    echo "Skipping home directory deletion."
;;

*)
    echo "Invalid option. Skipping home deletion."
;;
esac

echo
echo "====="
echo "User cleanup completed"
echo "====="

```

16.3 Make Script Executable

Before using the script:

```
chmod +x deleteUser.sh
```

16.4 Understanding Script Arguments

The script requires ONE argument:

Argument	Meaning
username	LDAP username

Syntax:

```
./deleteUser.sh <username>
```

Example:

```
./deleteUser.sh student02
```

16.5 What the Script Automatically Does

The script performs following operations automatically:

Operation	Description
Delete LDAP user	removes LDAP account
Optional group deletion	removes LDAP group if requested
Optional home deletion	removes user home if requested
Interactive confirmation	prevents accidental cleanup

16.6 Example User Deletion

Example:

```
./deleteUser.sh student02
```

Expected output:

```
=====
Deleting LDAP User
=====
Username : student02
=====
```

LDAP admin password will be requested.
After successful deletion:

LDAP user deleted successfully.

The script then asks:

Delete LDAP group 'student02'? (y/n):

If:

y

is selected, corresponding LDAP group will also be removed.
Next prompt:

Delete home directory '/home/student02'? (y/n):

If:

y

is selected, user home directory will also be deleted.

16.7 Important Group Deletion Notes

IMPORTANT:

Do NOT delete shared groups accidentally.

Example:

`hpcusers`

may contain multiple users.

Deleting shared groups can create:

- broken group mappings
- permission problems
- Slurm identity issues

The script therefore asks confirmation before deleting groups.

16.8 Important Home Directory Notes

Deleting home directories permanently removes:

- source codes
- job scripts
- outputs
- datasets
- SSH keys
- configuration files

IMPORTANT:

Home directory deletion is irreversible.

Always verify before deleting.

16.9 Verify User Removal

Verify deletion:

```
getent passwd student02
```

Expected:

```
no output
```

Verify LDAP entry removal:

```
ldapsearch -x \  
-b "dc=hpc,dc=local" uid=student02
```

Expected:

```
no entries returned
```

Verify login failure:

```
su - student02
```

Expected:

```
user does not exist
```

16.10 Important HPC User Management Concept

In centralized HPC environments:

- user lifecycle management
- account cleanup
- UID/GID consistency
- group management

are critical administrative tasks.

Improper deletion can create:

- orphaned files
- broken permissions
- Slurm scheduling issues
- incorrect ownership mappings

Therefore automated and controlled cleanup procedures are highly important in production HPC systems.

17 Automated LDAP Password Change Script

Managing passwords manually using LDIF files becomes inconvenient in centralized LDAP environments.

To simplify password management, an automated password change script can be used.

This script automates:

- LDAP password update
- password hash generation
- secure password input
- LDAP password modification

IMPORTANT:

Because authentication is centralized through LDAP, password changes automatically work cluster-wide.

After changing password:

- login node recognizes new password
- compute nodes recognize new password
- master recognizes new password

without changing anything individually on nodes.

17.1 Script Location

Example:

```
cd ~/ldap-scripts  
ls
```

Expected:

```
changePassword.sh
```


17.2 Password Change Script

Create file:

```
vim changePassword.sh
```

Paste the complete script exactly as shown below:

```
#!/bin/bash

LDAP_ADMIN="cn=Manager,dc=hpc,dc=local"
BASE_DN="dc=hpc,dc=local"

USERNAME=$1

if [ $# -ne 1 ]; then
    echo "Usage: $0 <username>"
    exit 1
fi

# Check if user exists
USER_EXISTS=$(ldapsearch -x \
-b "ou=People,$${BASE_DN}" "(uid=${USERNAME})" dn \
| grep "^dn:")

if [ -z "$USER_EXISTS" ]; then
    echo "ERROR: User ${USERNAME} does not exist!"
    exit 1
fi

echo
echo "====="
echo "LDAP Password Change"
echo "====="
echo "Username : ${USERNAME}"
echo "====="

# Take password securely
echo
read -s -p "Enter new password: " PASSWORD1
echo
```

```

read -s -p "Confirm new password: " PASSWORD2
echo

# Check password match
if [ "$PASSWORD1" != "$PASSWORD2" ]; then
    echo
    echo "ERROR: Passwords do not match!"
    exit 1
fi

# Generate LDAP password hash
PASSWORD_HASH=$(slappasswd -s "$PASSWORD1")

# Create LDIF file
cat > /tmp/${USERNAME}_password.ldif <<EOF
dn: uid=${USERNAME},ou=People,${BASE_DN}
changetype: modify
replace: userPassword
userPassword: ${PASSWORD_HASH}
EOF

echo
echo "Updating LDAP password..."

ldapmodify -x \
-D "${LDAP_ADMIN}" \
-W \
-f /tmp/${USERNAME}_password.ldif

if [ $? -ne 0 ]; then
    echo
    echo "ERROR: Failed to change password!"
    exit 1
fi

echo
echo "=====
echo "Password updated successfully"
echo "=====

```

```
echo "Username : ${USERNAME}"
echo "=====
```

17.3 Make Script Executable

Before using the script:

```
chmod +x changePassword.sh
```

17.4 Understanding Script Arguments

The script requires ONE argument:

Argument	Meaning
username	LDAP username

Syntax:

```
./changePassword.sh <username>
```

Example:

```
./changePassword.sh student01
```

17.5 What the Script Automatically Does

The script performs following operations automatically:

Operation	Description
Verify LDAP user exists	prevents invalid modifications
Secure password input	password hidden during typing
Password confirmation	avoids typing mistakes
Generate password hash	secure LDAP password storage
Modify LDAP password	centralized password update

17.6 Example Password Change

Example:

```
./changePassword.sh student01
```

Expected output:

```
=====
LDAP Password Change
=====
Username : student01
=====
```

The script asks:

Enter new password:

Then:

Confirm new password:

LDAP admin password will also be requested.

After successful update:

```
=====
Password updated successfully
=====
Username : student01
=====
```

17.7 Verify Password Change

Test login using updated password:

```
su - student01
```

or:

```
ssh student01@login
```

Expected:

```
successful login using new password
```

Because LDAP authentication is centralized, the updated password works automatically on:

- login node
- compute nodes
- master node

17.8 Important Password Management Notes

IMPORTANT:

LDAP stores password hashes, not plain-text passwords.

The script automatically generates secure LDAP-compatible password hashes using:

`slappasswd`

This improves security by avoiding direct plain-text password storage.